



**PES**  
**UNIVERSITY**

CELEBRATING **50** YEARS

# COMPUTER NETWORKS

---

**TEAM NETWORKS**

Department of Computer Science and Engineering

## Unit – 3 Network Layer and Internet Protocol

### 4.1 Overview of Network Layer

### 4.2 What's Inside a Router?

### 4.3 Switching

### 4.4 The Internet Protocol (IP)

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT

## Unit – 3 Network Layer and Internet Protocol

### 4.1 Overview of Network Layer

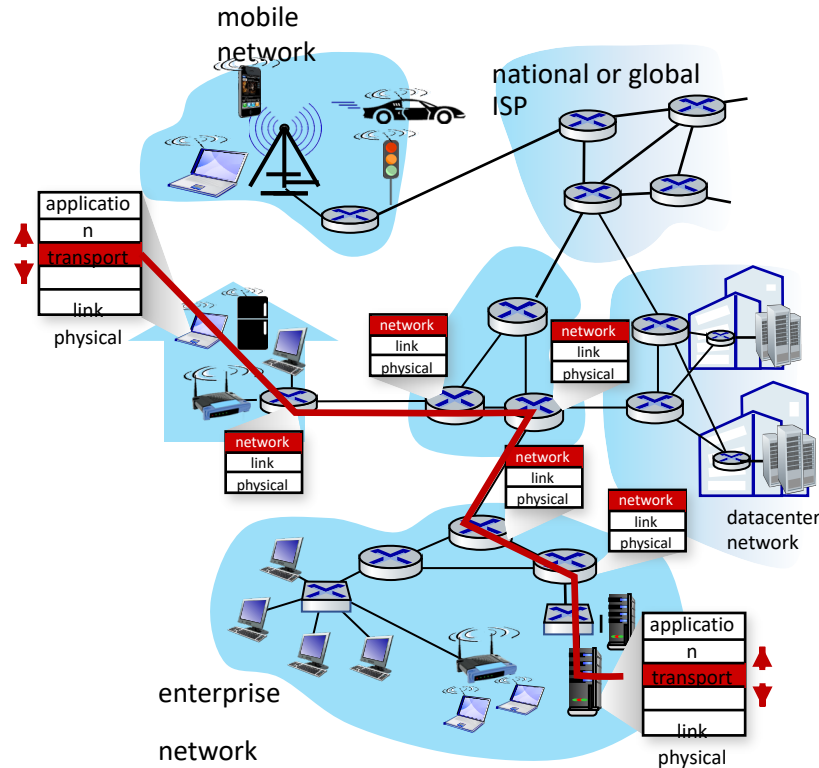
### 4.2 What's Inside a Router?

### 4.3 Switching

### 4.4 The Internet Protocol (IP)

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT

- transport segment from sending to receiving host
  - **sender:** encapsulates segments into datagrams, passes to link layer
  - **receiver:** delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- **routers:**
  - examines header fields in all IP datagrams passing through it
  - moves datagrams from input ports to output ports to transfer datagrams along end-end path



### network-layer functions:

- *forwarding*: move packets from a router's input link to appropriate router output link
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*

### analogy: taking a trip

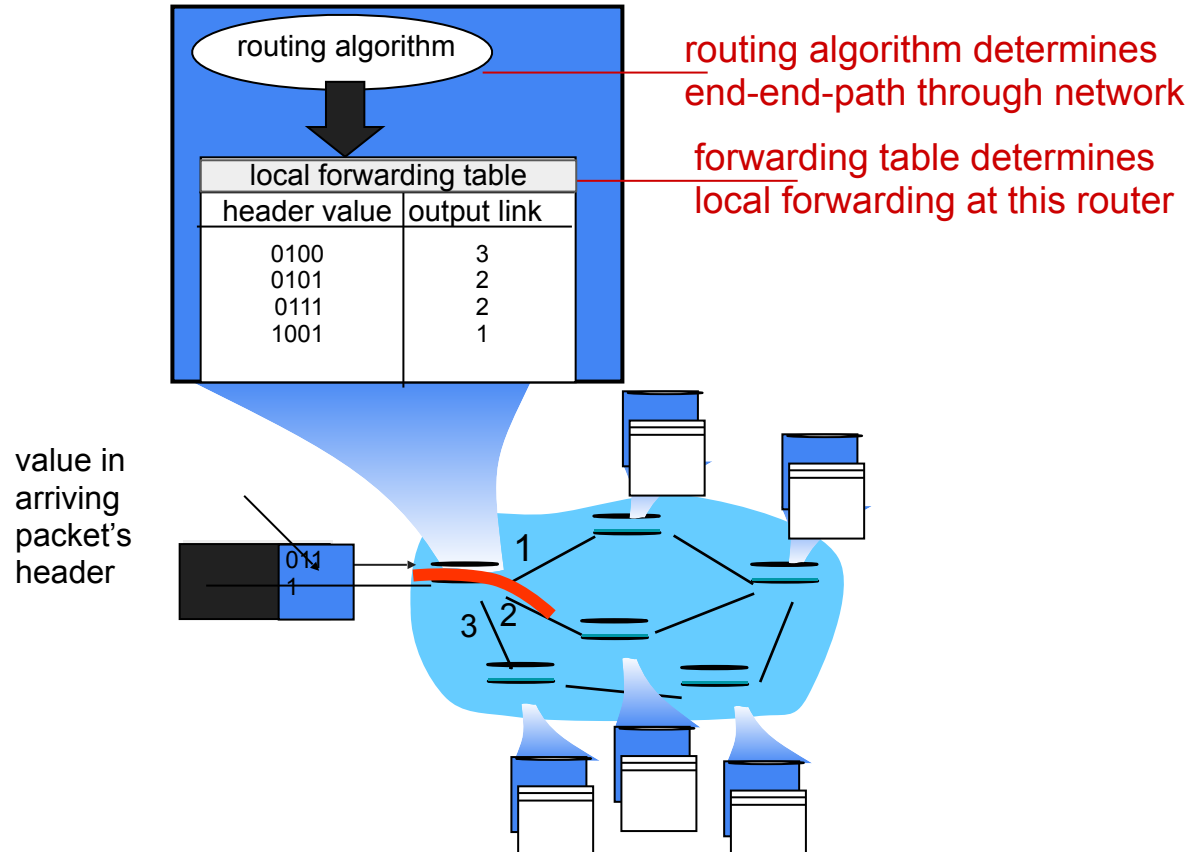
- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



forwarding



routing



- 3<sup>rd</sup> important function in *some* network architectures:
  - ATM, frame relay, X.25
- before datagrams flow, two end hosts *and* intervening routers establish virtual connection
  - routers get involved
- network vs transport layer connection service:
  - *network*: between two hosts (may also involve intervening routers in case of VCs)
  - *transport*: between two processes

**Q:** What *service model* for “channel” transporting datagrams from sender to receiver?

example services for *individual* datagrams:

- guaranteed delivery
- guaranteed delivery with less than 40 msec delay

example services for a *flow* of datagrams:

- in-order datagram delivery
- guaranteed minimum bandwidth to flow
- restrictions on changes in inter-packet spacing



| Network Architecture | Service Model | Quality of Service (QoS) Guarantees? |      |       |        |
|----------------------|---------------|--------------------------------------|------|-------|--------|
|                      |               | Bandwidth                            | Loss | Order | Timing |
| Internet             | best effort   | none                                 | no   | no    | no     |

Internet “best effort” service model

*No* guarantees on:

- successful datagram delivery to destination
- timing or order of delivery
- bandwidth available to end-end flow

| Network Architecture | Service Model                 | Quality of Service (QoS) Guarantees ? |          |          |        |
|----------------------|-------------------------------|---------------------------------------|----------|----------|--------|
|                      |                               | Bandwidth                             | Loss     | Order    | Timing |
| Internet             | best effort                   | none                                  | no       | no       | no     |
| ATM                  | Constant Bit Rate             | Constant rate                         | yes      | yes      | yes    |
| ATM                  | Available Bit Rate            | Guaranteed min                        | no       | yes      | no     |
| Internet             | Intserv Guaranteed (RFC 1633) | yes                                   | yes      | yes      | yes    |
| Internet             | Diffserv (RFC 2475)           | possible                              | possibly | possibly | no     |

## Unit – 4 Network Layer and Internet Protocol

### 4.1 Overview of Network Layer

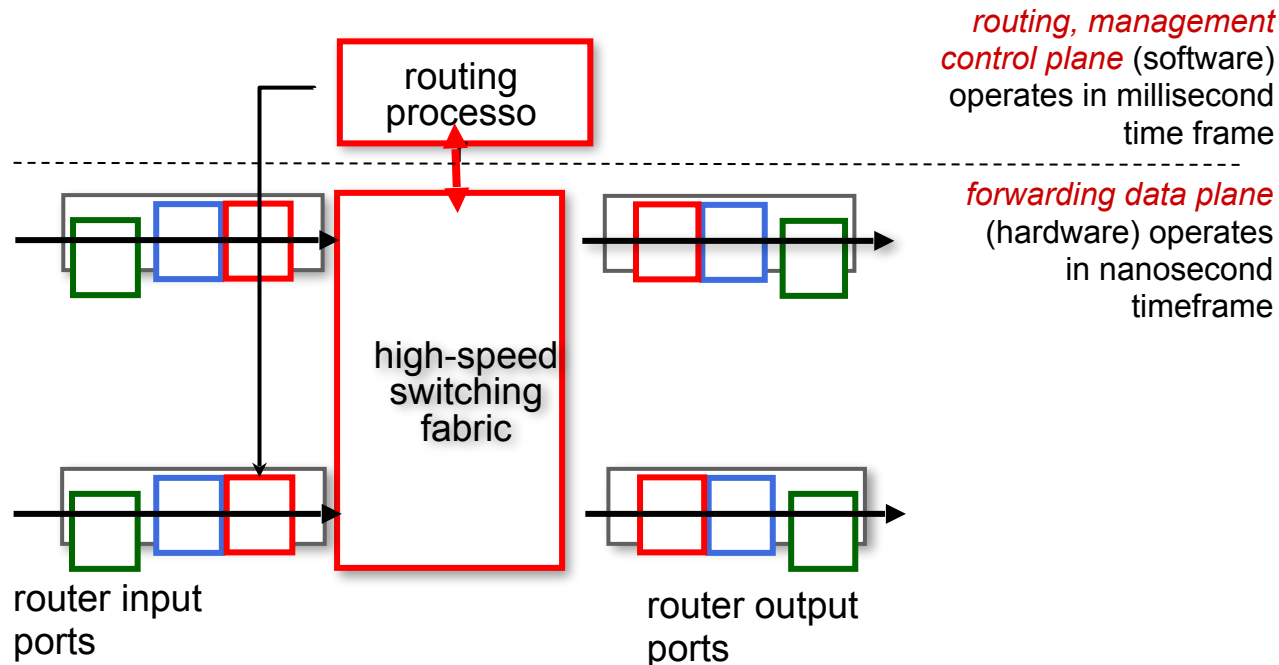
### 4.2 What's Inside a Router?

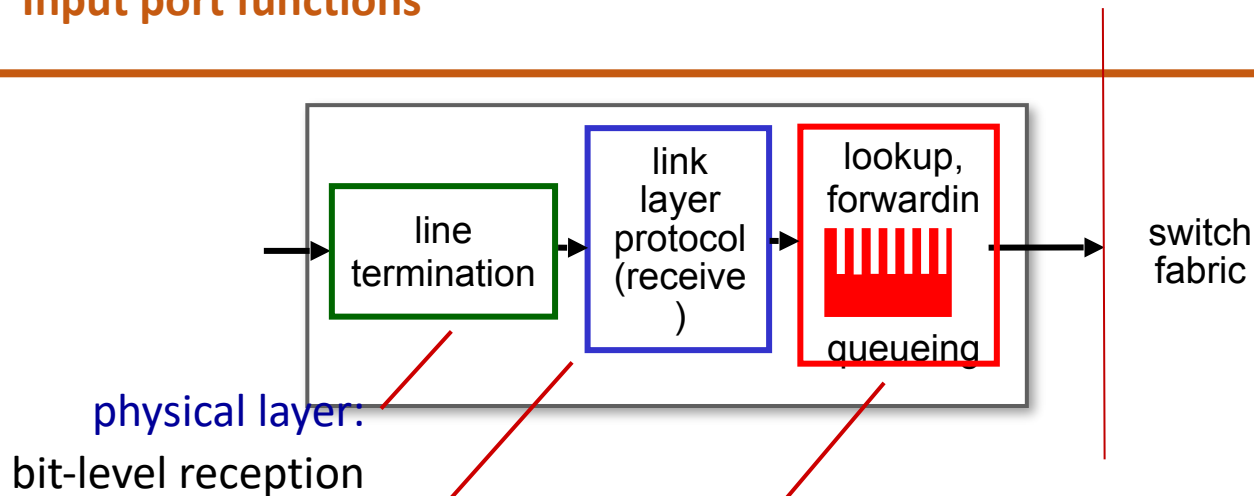
### 4.3 Switching

### 4.4 The Internet Protocol (IP)

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT

high-level view of generic router architecture:



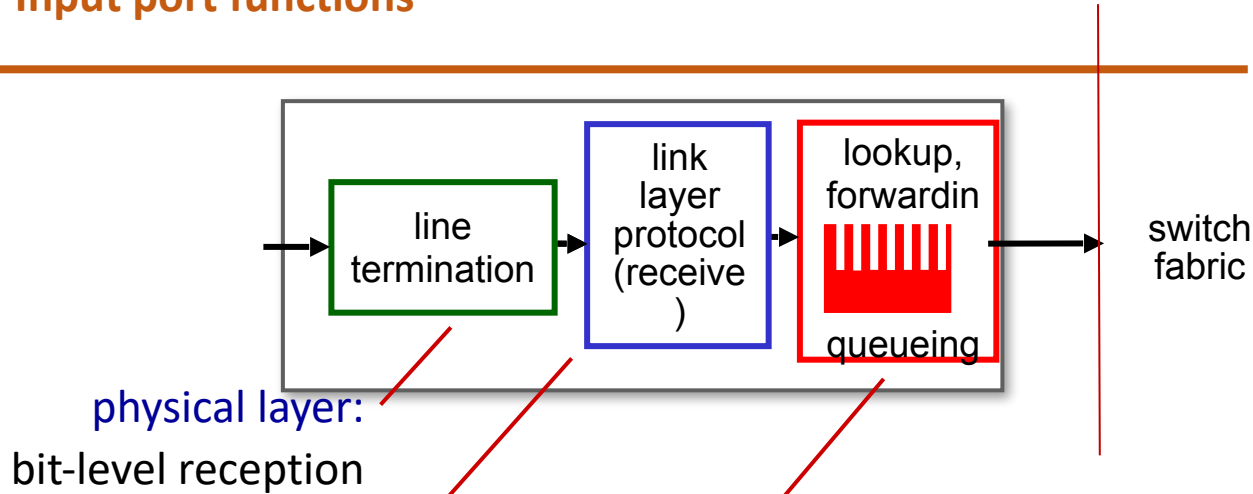


physical layer:  
bit-level reception

link layer:  
e.g., Ethernet  
(chapter 6)

### decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (*"match plus action"*)
- goal: complete input port processing at 'line speed'
- **input port queueing:** if datagrams arrive faster than forwarding rate into switch fabric



### decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (*"match plus action"*)
- **destination-based forwarding:** forward based only on destination IP address (traditional)
- **generalized forwarding:** forward based on any set of header field values

*forwarding table*

| Destination Address Range   | Link Interface |
|---|----------------|
| 11001000 00010111 00010000 00000000<br>through<br>11001000 00010111 00010000 00000100<br>through<br>11001000 00010111 00010000 00000111 | n<br>3         |
| 11001000 00010111 00011000 11111111   |                |
| 11001000 00010111 00011001 00000000<br>through<br>11001000 00010111 00011111 11111111   | 2              |
| otherwise   | 3              |

*Q:* but what happens if ranges don't divide up so nicely?

### longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range                                 | Link interface |
|---|----------------|
| 11001000    00010111    00010    * * *    * * * * * * * * | 0              |
| 11001000    00010111    00011000    * * * * * * * *       | 1              |
| 11001000    00010111    00011    * * *    * * * * * * * * | 2              |
| otherwise   | 3              |

examples:

11001000    00010111    00010110    10100001

which interface?

11001000    00010111    00011000    10101010

which interface?



### longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range  | Link interface |
|----------------------------|----------------|
| 11001000 00010111 00010*** | 0              |
| 11001000 00010111 00011000 | 1              |
| 11001000 00010111 00011000 | 2              |
| 11001000 00010111 00011*** | 3              |
| otherwise *                |                |

match!

examples

:

|                            |          |                  |
|----------------------------|----------|------------------|
| 11001000 00010111 00010110 | 10100001 | which interface? |
| 11001000 00010111 00011000 | 10101010 | which interface? |

### longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range        | Link      |
|----------------------------------|-----------|
| 11001000 00010111 00010* *****   | interface |
| 11001000 00010111 0001*000 ***** | 0         |
| 11001000 00010111 0001* *****    | 1         |
| otherwise *****                  | 2         |
|                                  | 3         |

match!

examples:

11001000 00010111 00010110  
 10100001  
 11001000 00010111 00011000  
 10101010

which interface?

which interface?

### longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range      | Link      |
|--------------------------------|-----------|
| 11001000 00010111 00010* * *   | interface |
| 11001000 00010111 0001100* * * | 0         |
| 11001000 00010111 00011* * *   | 1         |
| otherwise *                    | 2         |
|                                | 3         |

match!

examples:


11001000 00010111 00010110  
 10100001  
 11001000 00010111 0001100\* \* \*  
 10101010

which interface?




which interface?

- we'll see *why* longest prefix matching is used shortly, when we study addressing
- longest prefix matching: often performed using ternary content addressable memories (TCAMs)
  - *content addressable*: present address to TCAM: retrieve address in one clock cycle, regardless of table size
  - Cisco Catalyst: ~1M routing table entries in TCAM

- **switching rate:** rate at which packets can be transferred from inputs to outputs
  - often measured as multiple of input/output line rate
  - N inputs: switching rate N times line rate desirable




- transfer packet from input link to appropriate output link
- **switching rate**: rate at which packets can be transferred from inputs to outputs
  - often measured as multiple of input/output line rate
  - N inputs: switching rate N times line rate desirable
- three major types of switching fabrics:




### first generation routers:

- traditional computers with switching under direct control of CPU
- packet copied to system's memory
- speed limited by memory bandwidth (2 bus crossings per datagram)




- datagram from input port memory to output port memory via a shared bus
- *bus contention*: switching speed limited by bus bandwidth
- 32 Gbps bus, Cisco 5600: sufficient speed for access routers






- Crossbar, Clos networks, other interconnection nets initially developed to connect processors in multiprocessor
- **multistage switch**:  $n \times n$  switch from multiple stages of smaller switches
- **exploiting parallelism**:
  - fragment datagram into fixed length cells on entry
  - switch cells through the fabric, reassemble datagram at exit




3x3  
crossbar





8x8 multistage switch  
built from smaller-sized  
switches

- scaling, using multiple switching “planes” in parallel:
  - speedup, scaleup via parallelism
- Cisco CRS router:
  - basic unit: 8 switching planes
  - each plane: 3-stage interconnection network
  - up to 100's Tbps switching capacity




- If switch fabric slower than input ports combined -> queueing may occur at input queues
  - queueing delay and loss due to input buffer overflow!
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward





This is a really important slide



- **Buffering** required when datagrams arrive from fabric faster than link transmission rate. **Drop policy:** which datagrams to drop if no free buffers?

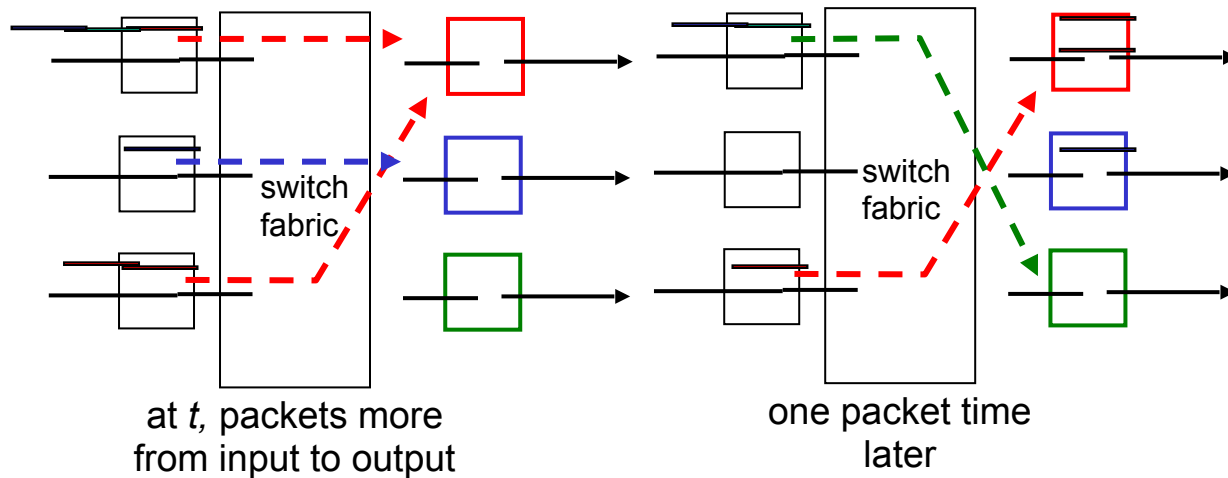


Datagrams can be lost due to congestion, lack of buffers

- **Scheduling discipline** chooses among queued datagrams for transmission



Priority scheduling – who gets best performance, network neutrality

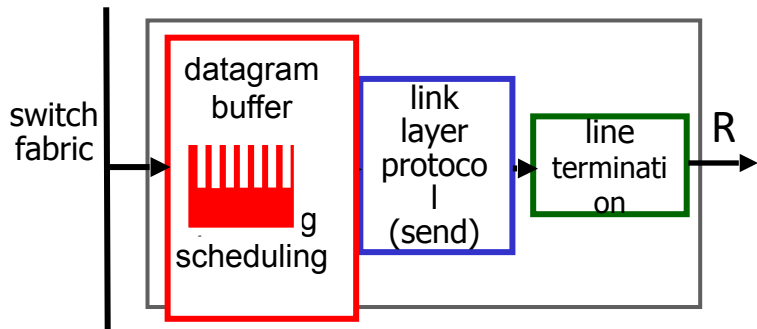


- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to output port buffer overflow!*

- RFC 3439 rule of thumb: average buffering equal to “typical” RTT (say 250 msec) times link capacity  $C$ 
  - e.g.,  $C = 10$  Gbps link: 2.5 Gbit buffer
- more recent recommendation: with  $N$  flows, buffering equal to

$$\frac{RTT \cdot C}{\sqrt{N}}$$


- but *too* much buffering can increase delays (particularly in home routers)
  - long RTTs: poor performance for realtime apps, sluggish TCP response
  - recall delay-based congestion control: “keep bottleneck link just full enough (busy) but no fuller”



### buffer management:

- **drop**: which packet to add, drop when buffers are full
  - **tail drop**: drop arriving packet
  - **priority**: drop/remove on priority basis

### Abstraction: queue



- **marking**: which packets to mark to signal congestion (ECN, RED)

## Unit – 4 Network Layer and Internet Protocol

4.1 Overview of Network Layer

4.2 What's Inside a Router?

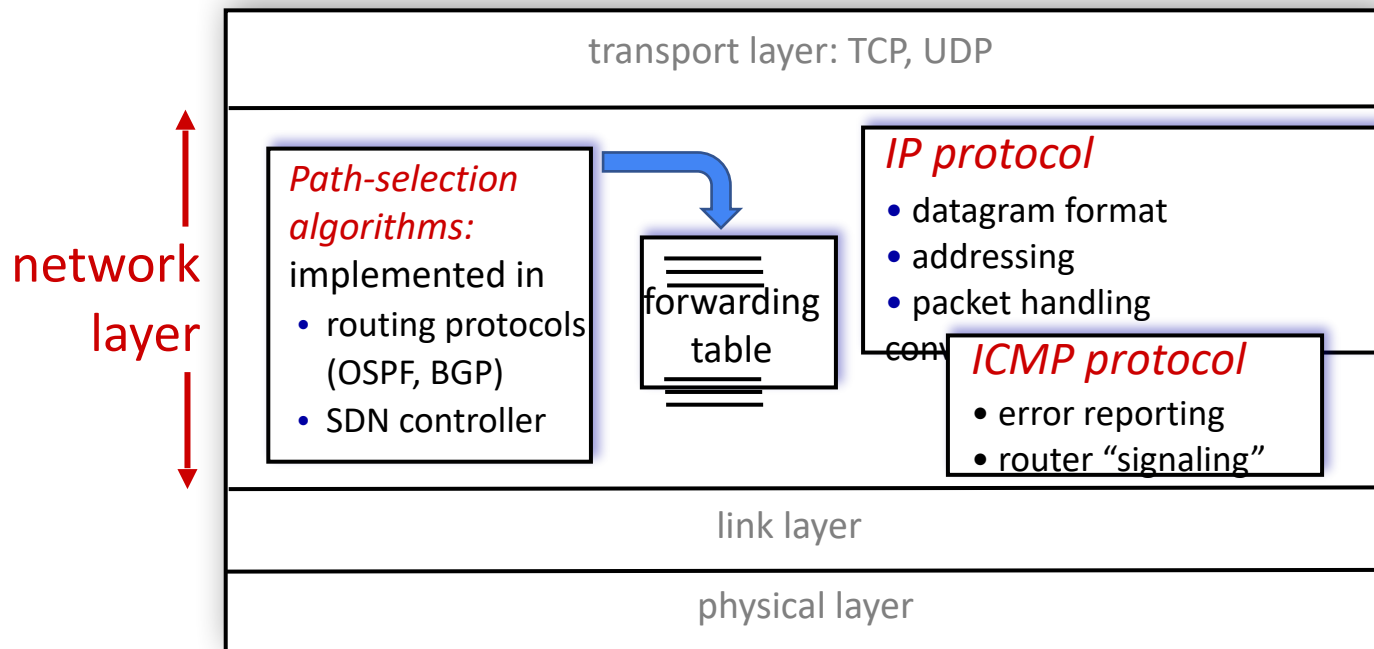
4.3 Switching

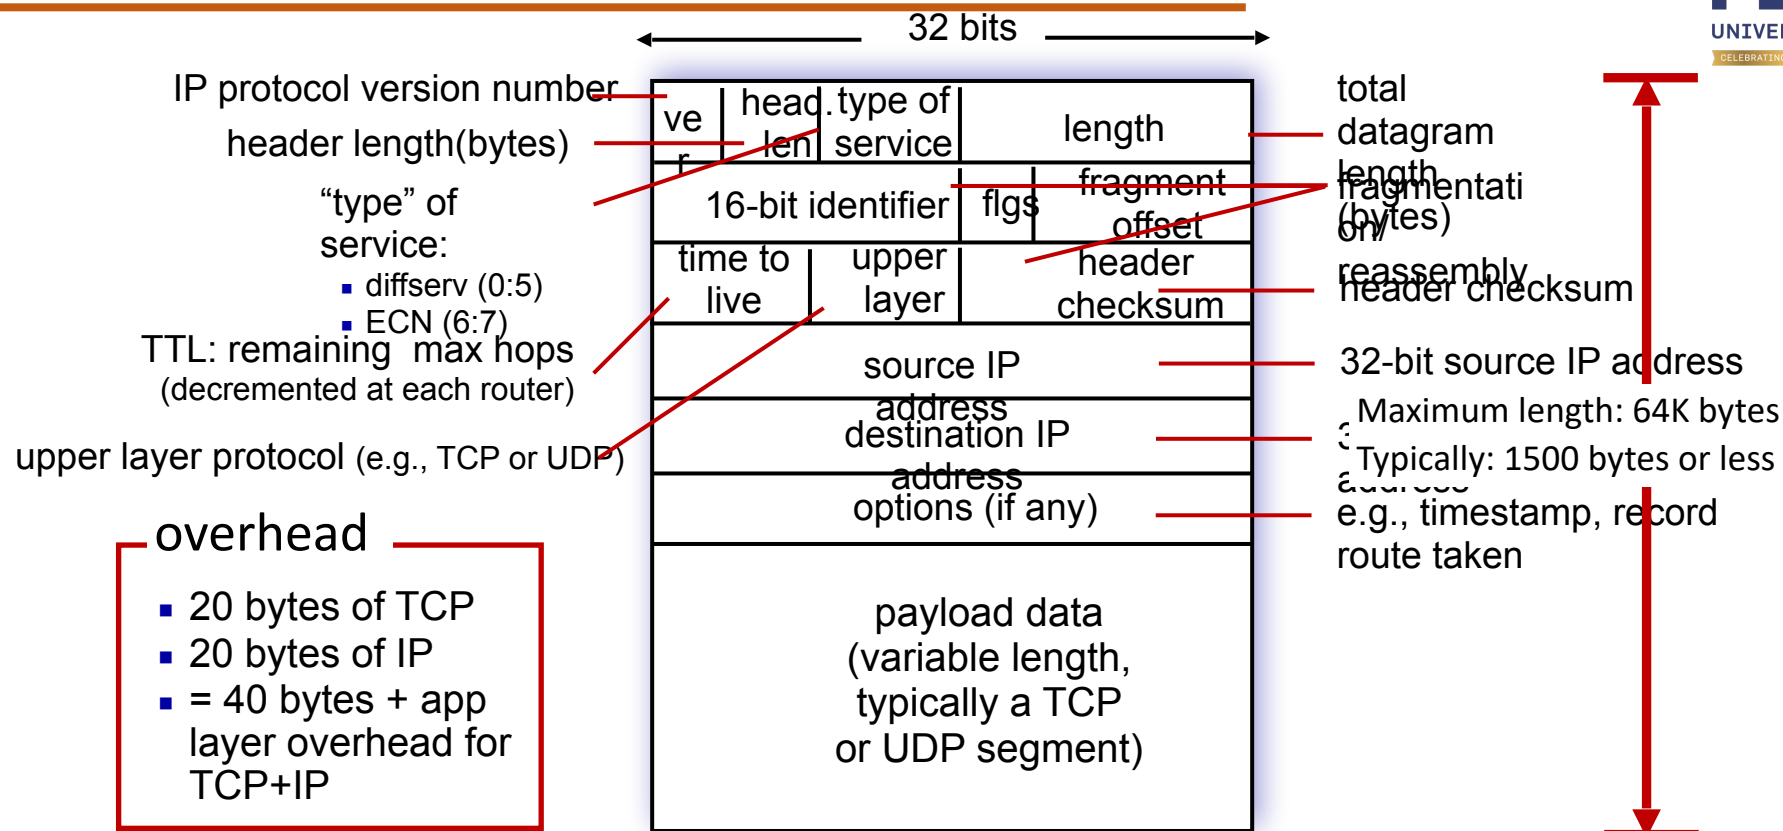
**4.4 The Internet Protocol (IP)**

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT



host, router network layer functions:





## Unit – 4 Network Layer and Internet Protocol

4.1 Overview of Network Layer


4.2 What's Inside a Router?

4.3 Switching

**4.4 The Internet Protocol (IP)**

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT

- network links have MTU (max.transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - “reassembled” only at final destination
  - IP header bits used to identify, order related fragments



*example:*

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

|  |                 |          |                |              |  |
|--|-----------------|----------|----------------|--------------|--|
|  | length<br>=4000 | ID<br>=x | fragflag<br>=0 | offset<br>=0 |  |
|--|-----------------|----------|----------------|--------------|--|

*one large datagram becomes  
several smaller datagrams*

1480 bytes  
in data field

|  |                 |          |                |              |  |
|--|-----------------|----------|----------------|--------------|--|
|  | length<br>=1500 | ID<br>=x | fragflag<br>=1 | offset<br>=0 |  |
|--|-----------------|----------|----------------|--------------|--|

|  |                 |          |                |                |  |
|--|-----------------|----------|----------------|----------------|--|
|  | length<br>=1500 | ID<br>=x | fragflag<br>=1 | offset<br>=185 |  |
|--|-----------------|----------|----------------|----------------|--|

offset =  $1480/8$

|  |                 |          |                |                |  |
|--|-----------------|----------|----------------|----------------|--|
|  | length<br>=1040 | ID<br>=x | fragflag<br>=0 | offset<br>=370 |  |
|--|-----------------|----------|----------------|----------------|--|

### Original IP Datagram

| Sequence | Identifier | Total Length | DF<br>May / Don't | MF<br>Last / More | Fragment<br>Offset |
|----------|------------|--------------|-------------------|-------------------|--------------------|
| 0        | 345        | 5140         | 0                 | 0                 | 0                  |

### IP Fragments (Ethernet)

| Sequence | Identifier | Total Length | DF<br>May / Don't | MF<br>Last / More | Fragment<br>Offset |
|----------|------------|--------------|-------------------|-------------------|--------------------|
| 0-0      | 345        | 1500         | 0                 | 1                 | 0                  |
| 0-1      | 345        | 1500         | 0                 | 1                 | 185                |
| 0-2      | 345        | 1500         | 0                 | 1                 | 370                |
| 0-3      | 345        | 700          | 0                 | 0                 | 555                |

## Fragmentation – Numerical Example

- An IP router with a Maximum Transmission Unit (MTU) of 200 bytes has received an IP packet of size 520 bytes with an IP header of length 20 bytes. The values of the relevant fields in the IP header.

|                 | <div>20</div> <div>176</div> | <div>20</div> <div>176</div> | <div>20</div> <div>148</div> |
|-----------------|------------------------------|------------------------------|------------------------------|
| Fragment Offset | 0                            | 22                           | 44                           |
| MF              | 1                            | 1                            | 0                            |
| Header length   | 5                            | 5                            | 5                            |
| Total length    | 196                          | 196                          | 168                          |



## Unit – 4 Network Layer and Internet Protocol

4.1 Overview of Network Layer

4.2 What's Inside a Router?


4.3 Switching

### 4.4 The Internet Protocol (IP)

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT



- **IP address:** 32-bit unique ID associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



dotted-decimal IP address notation:

223.1.1.1 = 11011111 00000001 00000001 00000001  
223                    1                    1                    1

Change the following IPv4 addresses from binary notation to dotted-decimal notation.

a. 10000001 00001011 00001011 11101111

b. 11000001 10000011 00011011 11111111

### Solution:

We replace each group of 8 bits with its equivalent decimal number and add dots for separation.

a. 129.11.11.239

b. 193.131.27.255



Change the following IPv4 addresses from dotted-decimal notation to binary notation.

a. 111.56.45.78

b. 221.34.7.82

### Solution:

We replace each decimal number with its binary equivalent.

a. 01101111 00111000 00101101 01001110

b. 11011101 00100010 00000111 01010010




**Q:** how are interfaces  
actually connected?

**A:** we'll learn about that in  
chapters 6, 7

**A:** wired  
Ethernet interfaces  
connected by  
Ethernet switches

**For now:** don't need to worry  
about how one interface is  
connected to another (with no  
intervening router)



**A:** wireless WiFi interfaces  
connected by WiFi base station


### ■ *What's a subnet ?*

- device interfaces that can physically reach each other **without passing through an intervening router**

### ■ IP addresses have structure:

- **subnet part:** devices in same subnet have common high order bits
- **host part:** remaining low order bits


**A portion of a network that shares a particular subnet address.**



network consisting of 3 subnets

### *Recipe for defining subnets:*


- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a *subnet*



subnet mask: /24

(high-order 24 bits: subnet part of IP address)


- where are the subnets?
- what are the /24 subnet addresses?



A 32-bit combination used to describe which portion of an address refers to the subnet and which part refers to the host.

### CIDR: Classless InterDomain Routing (pronounced “cider”)

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address





# COMPUTER NETWORKS

## IP addressing: CIDR

| CIDR Block Prefix | # of Host Addresses (inclusive of network and broadcast-id) |
|-------------------|---|
|-------------------|---|

|     |                                   |
|-----|-----------------------------------|
| /30 | 4 hosts (valid host would be 2)   |
| /29 | 8 hosts (valid host would be 6)   |
| /28 | 16 hosts (valid host would be 14) |
| /27 | 32 hosts                          |
| /26 | 64 hosts                          |
| /25 | 128 hosts                         |
| /24 | 256 hosts                         |
| /23 | 512 hosts                         |
| /22 | 1,024 hosts                       |
| /21 | 2,048 hosts                       |
| /20 | 4,096 hosts                       |
| /19 | 8,192 hosts                       |
| /18 | 16,384 hosts                      |
| /17 | 32,768 hosts                      |
| /16 | 65,536 hosts                      |
| /15 | 131,072 hosts                     |
| /14 | 262,144 hosts                     |
| /13 | 524,288 hosts                     |
| /12 | 1,048,576 hosts and so on.....    |

- Network portion - 8, 16, or 24 bits in length – known as **Class A, B and C** networks respectively
- A class **C (/24)** - accommodate only up to  $2^8 - 2 = 254$  **hosts** (two of the  $2^8 = 256$  addresses are reserved for special use)
- A class **B (/16)** subnet, supports up to **65,634 hosts**
- A class **A (/8)** subnet, up to **16,777,214 hosts**

**too small for many organizations**

**too large, poor utilization**

1

$2^n = \text{Number of Network}$






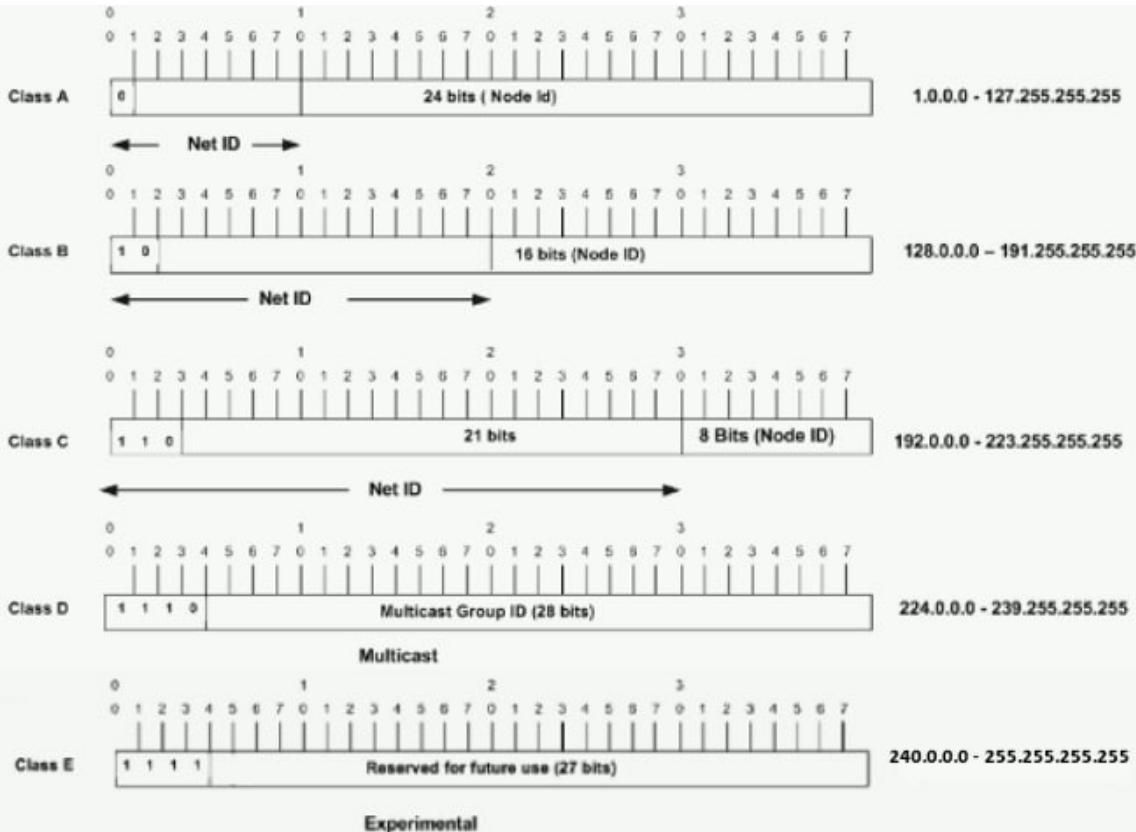
2

$2^h - 2 = \text{Number of Host}$

**Classful addressing, which is almost obsolete, is replaced with classless addressing.**

# COMPUTER NETWORKS

## IP addressing: Classful Addressing



# COMPUTER NETWORKS

## IP addressing: Classful Addressing

| Class   | HOB  | NET ID Bits | Host ID Bits | No of Networks     | Host Per Network    | Start Address | End Address     |
|---------|------|-------------|--------------|--------------------|---------------------|---------------|-----------------|
| Class A | 0    | 8           | 24           | $2^7=128$          | $2^{24}=16,777,216$ | 0.0.0.0       | 127.255.255.255 |
| Class B | 10   | 16          | 16           | $2^{14}=16,384$    | $2^{16}=65,536$     | 128.0.0.0     | 191.255.255.255 |
| Class C | 110  | 24          | 8            | $2^{21}=2,097,152$ | $2^8=256$           | 192.0.0.0     | 223.255.255.255 |
| Class D | 1110 | -           | -            | -                  | -                   | 224.0.0.0     | 239.255.255.255 |
| Class E | 1111 | -           | -            | -                  | -                   | 240.0.0.0     | 255.255.255.255 |

**Class D: multicast, Class E: reserved (experimental)**

**Broadcast Address: 255.255.255.255**

### Network Masks

Class A: 255.0.0.0

Class B: 255.255.0.0

Class C: 255.255.255.0

# COMPUTER NETWORKS

## IP addressing: Classful Subnet Masks

| Class    | Binary                              | Dotted-Decimal | CIDR Notation |
|----------|-------------------------------------|----------------|---------------|
| <b>A</b> | 11111111 00000000 00000000 00000000 | 255.0.0.0      | /8            |
| <b>B</b> | 11111111 11111111 00000000 00000000 | 255.255.0.0    | /16           |
| <b>C</b> | 11111111 11111111 11111111 00000000 | 255.255.255.0  | /24           |

**Q:** how does *network* get subnet part of IP address?

**A:** gets allocated portion of its provider ISP's address space

ISP's block      11001000 00010111 00010000 00000000    200.23.16.0/20

ISP can then allocate out its address space in 8 blocks:

|                |                                   |          |                |
|----------------|-----------------------------------|----------|----------------|
| Organization 0 | <u>11001000 00010111 00010000</u> | 00000000 | 200.23.16.0/23 |
| Organization 1 | <u>11001000 00010111 00010010</u> | 00000000 | 200.23.18.0/23 |
| Organization 2 | <u>11001000 00010111 00010100</u> | 00000000 | 200.23.20.0/23 |
| ...            | .....                             | ....     | ....           |
| Organization 7 | <u>11001000 00010111 00011110</u> | 00000000 | 200.23.30.0/23 |

**Find the class of each address.**

- a. 00000001 00001011 00001011 11101111
- b. 11000001 10000011 00011011 11111111
- c. 14.23.120.8
- d. 252.5.15.111

**Solution**

- a. The first bit is 0. This is a class A address.
- b. The first 2 bits are 1; the third bit is 0. This is a class C address.
- c. The first byte is 14; the class is A.
- d. The first byte is 252; the class is E.



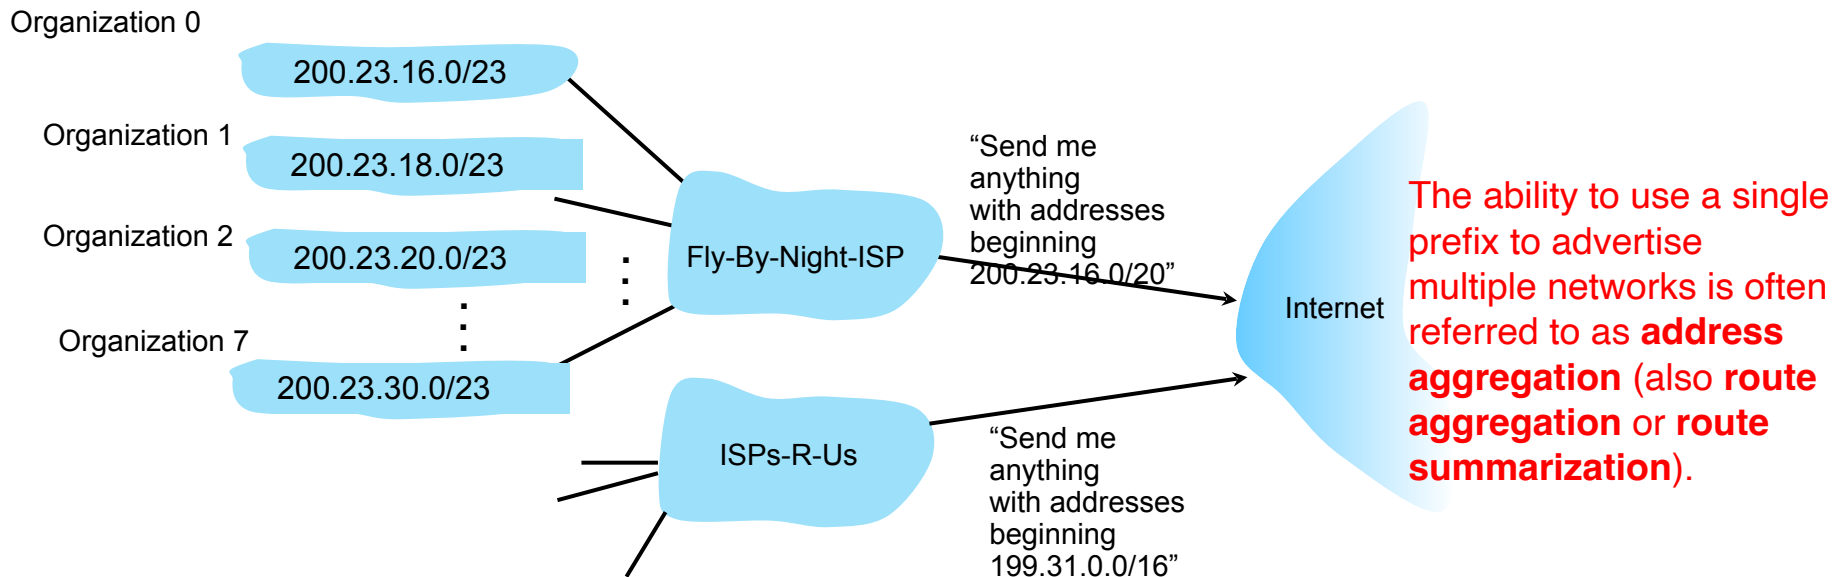
| <i>/n</i> | <i>Mask</i> | <i>/n</i> | <i>Mask</i> | <i>/n</i> | <i>Mask</i>   | <i>/n</i> | <i>Mask</i>     |
|-----------|-------------|-----------|-------------|-----------|---------------|-----------|-----------------|
| /1        | 128.0.0.0   | /9        | 255.128.0.0 | /17       | 255.255.128.0 | /25       | 255.255.255.128 |
| /2        | 192.0.0.0   | /10       | 255.192.0.0 | /18       | 255.255.192.0 | /26       | 255.255.255.192 |
| /3        | 224.0.0.0   | /11       | 255.224.0.0 | /19       | 255.255.224.0 | /27       | 255.255.255.224 |
| /4        | 240.0.0.0   | /12       | 255.240.0.0 | /20       | 255.255.240.0 | /28       | 255.255.255.240 |
| /5        | 248.0.0.0   | /13       | 255.248.0.0 | /21       | 255.255.248.0 | /29       | 255.255.255.248 |
| /6        | 252.0.0.0   | /14       | 255.252.0.0 | /22       | 255.255.252.0 | /30       | 255.255.255.252 |
| /7        | 254.0.0.0   | /15       | 255.254.0.0 | /23       | 255.255.254.0 | /31       | 255.255.255.254 |
| /8        | 255.0.0.0   | /16       | 255.255.0.0 | /24       | 255.255.255.0 | /32       | 255.255.255.255 |

The addresses in color are the default masks for classes A, B, and C.

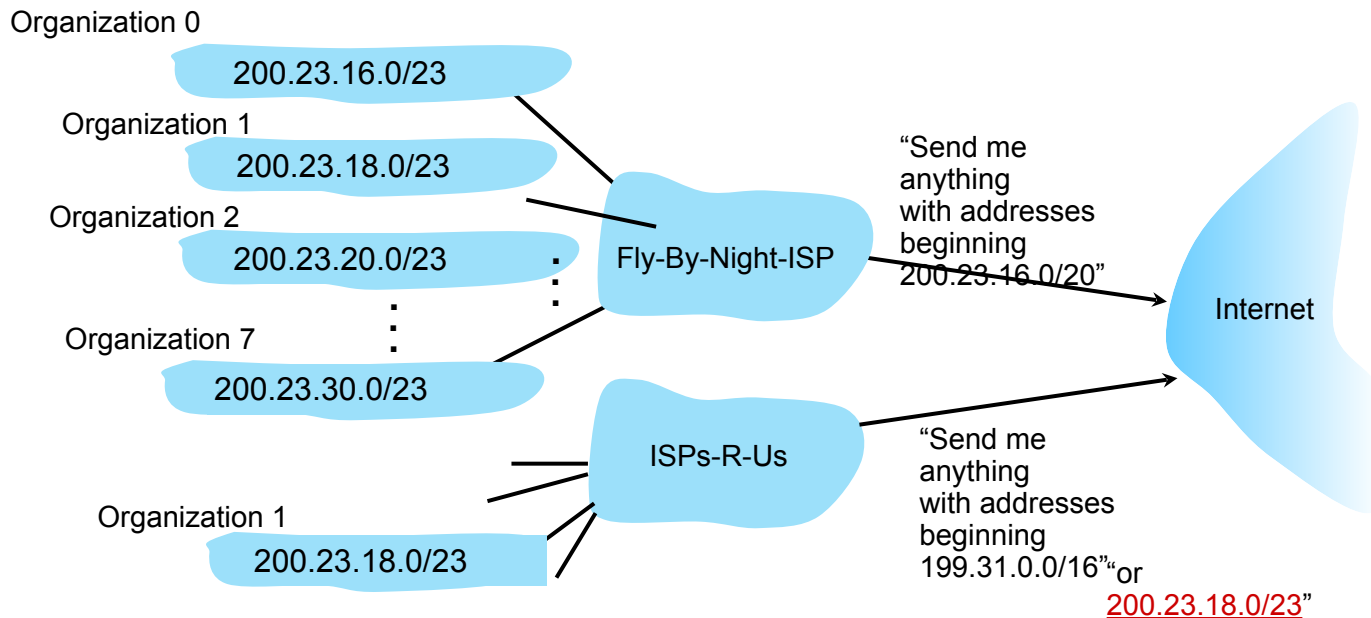
Thus, classful addressing is a special case of classless addressing.



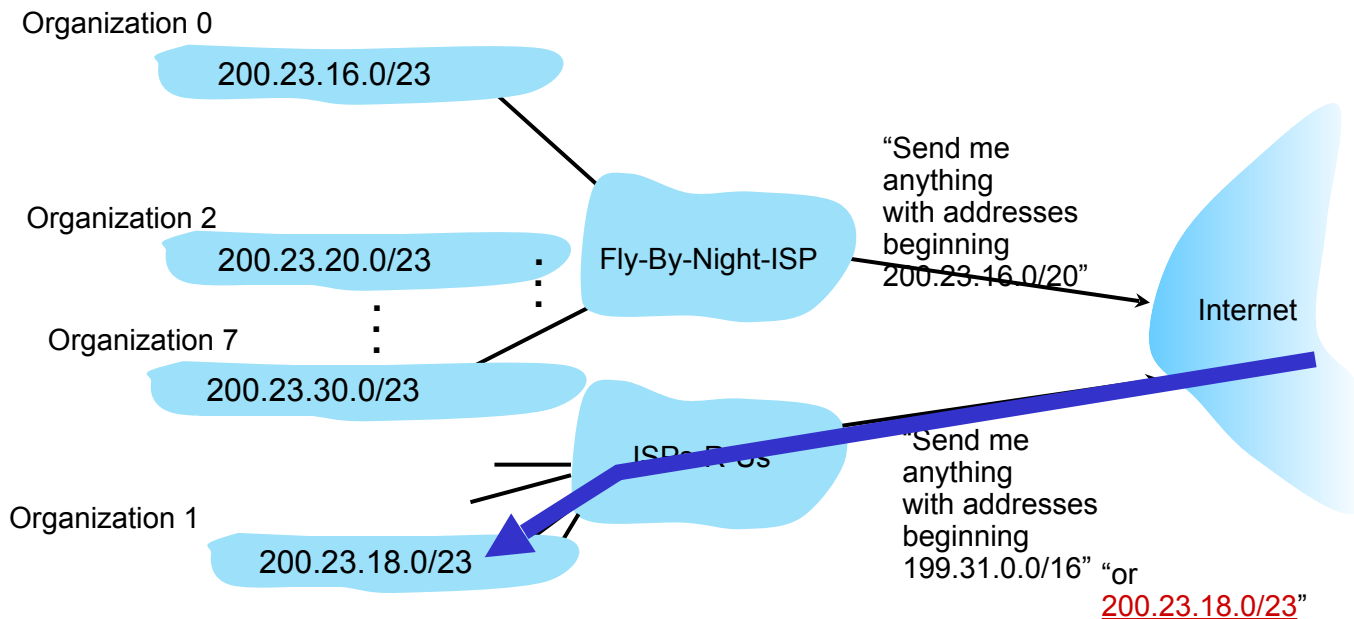
hierarchical addressing allows efficient advertisement of routing information:



- Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



- Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



**Q:** how does an ISP get block of addresses?

**A: ICANN:** Internet Corporation for Assigned Names and Numbers <http://www.icann.org/>

- allocates IP addresses, through 5 regional registries (RRs) (who may then allocate to local registries)
- manages DNS root zone, including delegation of individual TLD (.com, .edu, ...) management

**Q:** are there enough 32-bit IP addresses?

- ICANN allocated last chunk of IPv4 addresses to RRs in 2011
- NAT (next) helps IPv4 address space exhaustion
- IPv6 has 128-bit address space

"Who the hell knew how much address space we needed?" Vint Cerf (reflecting on decision to make IPv4 address 32 bits long)

## Unit – 4 Network Layer and Internet Protocol

4.1 Overview of Network Layer

4.2 What's Inside a Router?


4.3 Switching

### 4.4 The Internet Protocol (IP)

- Datagram format
- Fragmentation
- IPv4 addressing
- NAT

- all devices in local network have 32-bit addresses in a “private” IP address space (10/8, 172.16/12, 192.168/16 prefixes) that can only be used in local network
- advantages:
  - just **one** IP address needed from provider ISP for *all* devices
  - can change addresses of host in local network without notifying outside world
  - can change ISP without changing addresses of devices in local network
  - security: devices inside local net not directly addressable, visible by outside world

**NAT:** all devices in local network share just **one** IPv4 address as far as outside world is concerned



*all* datagrams *leaving* local network have *same* source NAT IP address: 138.76.29.7, but *different* source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

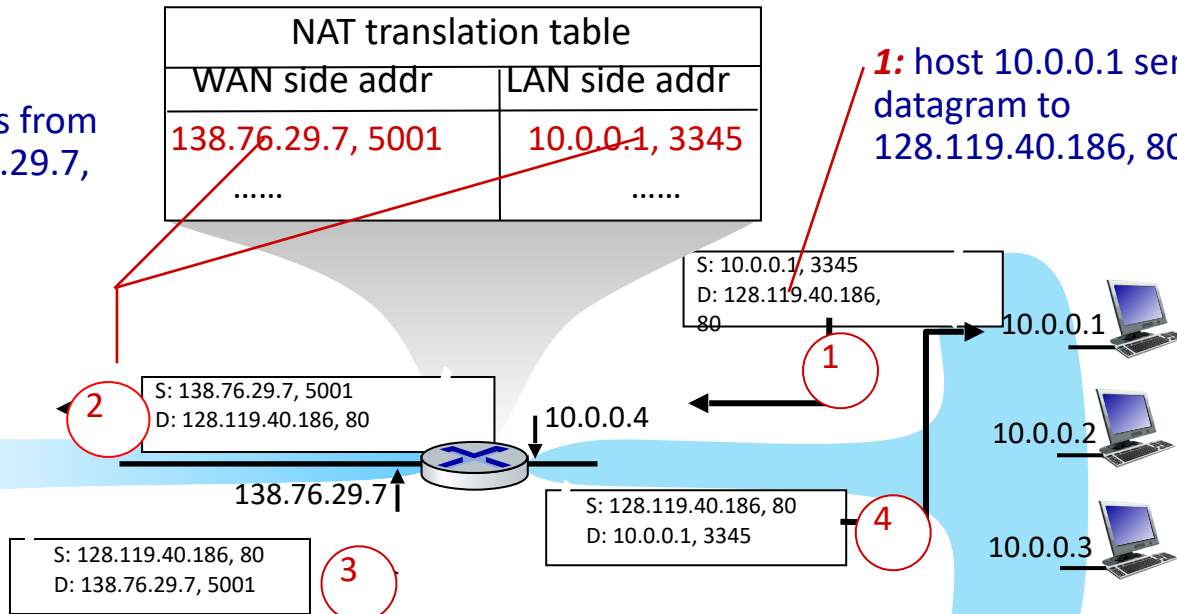
**implementation:** NAT router must (transparently):

- **outgoing datagrams: replace** (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
  - remote clients/servers will respond using (NAT IP address, new port #) as destination address
- **remember (in NAT translation table)** every (source IP address, port #) to (NAT IP address, new port #) translation pair
- **incoming datagrams: replace** (NAT IP address, new port #) in destination fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table




**2:** NAT router changes datagram source address from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table

**1:** host 10.0.0.1 sends datagram to 128.119.40.186, 80



**3:** reply arrives, destination address: 138.76.29.7, 5001

- NAT has been controversial:
  - routers “should” only process up to layer 3
  - address “shortage” should be solved by IPv6
  - violates end-to-end argument (port # manipulation by network-layer device)
  - NAT traversal: what if client wants to connect to server behind NAT?
- but NAT is here to stay:
  - extensively used in home and institutional nets, 4G/5G cellular nets



***Solution:*** statically configure NAT  
to forward incoming connection  
requests at given port to server  
e.g., (123.76.29.7, port 25000)  
always forwarded to 10.0.0.1 port  
25000

Find the range of addresses in the following blocks.

- 123.56.77.32/29
- 180.34.64.64/30
- 200.17.21.128/27



You are given the network address **175.200.0.0**; you are required to have 4 subnets.

What is the minimum number of Host Bits can you take in to the Network Bits for this purpose?

Write down the addresses of 4 subnets.

Write the subnet mask for the network.



In a block of addresses, we know the IP address of one host is **25.34.12.56/16**.

What are the first address (network address) and the last address (limited broadcast address) in this block?

Please explain in full with detailed steps showing the calculation.



An organization is granted the block **16.0.0.0/8**. The administrator wants to create 500 fixed-length subnets.

- i) find the subnet mask
- ii) find the number of addresses in each subnet.
- iii) find the first and last address in first subnet and in the last subnet (500<sup>th</sup> subnet).



An organization is granted a block of addresses with the beginning address **14.24.74.0/24**. The organization needs to have 3 subblocks of addresses to use in its three subnets: one subblock of **10 addresses**, one subblock of **60 addresses**, and one subblock of **120 addresses**. Design the subblocks.





Find the network address and the directed broadcast address of subnetted Class B IPv4 address **172.25.171.182** with a subnet mask of **255.255.224.0**.





**THANK YOU**

---

**TEAM NETWORKS**

Department of Computer Science and Engineering